

Before you turn this problem in, make sure everything runs as expected. First, **restart the kernel** (in the menu bar, select Kernel→Restart) and then **run all cells** (in the menu bar, select Cell→Run All).

Make sure that in addition to the code, you provide written answers for all questions of the assignment.

Below, please fill in your name:

```
In [1]: NAME = "Jesus Andres Correal Ortiz"
```

Assignment 1 - Data Cleaning & Visualization

(30 points total)

For this assignment, we will use an open dataset with data on various types of cereal and the corresponding customer ratings. Use the .csv file provided.

Create a dataframe by importing the file into Jupyter, and complete the tasks below. Be sure to show your Python code. You will not be graded on code efficiency, but your code should return a correct answer.

Dictionary

Content

Fields in the dataset:

- Name: Name of cereal
- mfr: Manufacturer of cereal
 - o A = American Home Food Products;
 - o G = General Mills
 - o K = Kelloggs
 - o N = Nabisco
 - o P = Post
 - o Q = Quaker Oats
 - o R = Ralston Purina
- type:
 - o cold
 - o hot
- calories: calories per serving
- protein: grams of protein
- fat: grams of fat
- sodium: milligrams of sodium
- fiber: grams of dietary fiber
- carbo: grams of complex carbohydrates
- sugars: grams of sugars
- potass: milligrams of potassium
- vitamins: vitamins and minerals – 0, 25, or 100, indicating the typical percentage of FDA recommended

- shelf: display shelf (1, 2, or 3, counting from the floor)
- weight: weight in ounces of one serving
- cups: number of cups in one serving
- rating: a rating of the cereals

Requirements

```
In [2]: %pip install pandas==2.1.2
        %pip install matplotlib==3.8.0
        %pip install numpy==1.26.1
```

Collecting pandas==2.1.2

Obtaining dependency information for pandas==2.1.2 from https://files.pythonhosted.org/packages/4e/dd/4a77fb4cb7d207fbeb77dfc7c022131d295767504eabb5836fcd63b644a1/pandas-2.1.2-cp311-cp311-macosx_11_0_arm64.whl.metadata

Downloading pandas-2.1.2-cp311-cp311-macosx_11_0_arm64.whl.metadata (18 kB)
Requirement already satisfied: numpy<2,>=1.23.2 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from pandas==2.1.2) (1.24.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from pandas==2.1.2) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from pandas==2.1.2) (2022.7)
Collecting tzdata>=2022.1 (from pandas==2.1.2)

Using cached tzdata-2023.3-py2.py3-none-any.whl (341 kB)
Requirement already satisfied: six>=1.5 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from python-dateutil>=2.8.2->pandas==2.1.2) (1.16.0)

Using cached pandas-2.1.2-cp311-cp311-macosx_11_0_arm64.whl (10.8 MB)

Installing collected packages: tzdata, pandas

Attempting uninstall: pandas

Found existing installation: pandas 1.5.3

Uninstalling pandas-1.5.3:

Successfully uninstalled pandas-1.5.3

Successfully installed pandas-2.1.2 tzdata-2023.3

Note: you may need to restart the kernel to use updated packages.

Collecting matplotlib==3.8.0

Obtaining dependency information for matplotlib==3.8.0 from https://files.pythonhosted.org/packages/af/f3/fb27b3b902fc759bbca3f9d0336c48069c3022e57552c4b0095d997c7ea8/matplotlib-3.8.0-cp311-cp311-macosx_11_0_arm64.whl.metadata

Downloading matplotlib-3.8.0-cp311-cp311-macosx_11_0_arm64.whl.metadata (5.8 kB)

Requirement already satisfied: contourpy>=1.0.1 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (1.0.5)
Requirement already satisfied: cycler>=0.10 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (1.4.4)
Requirement already satisfied: numpy<2,>=1.21 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (1.24.3)
Requirement already satisfied: packaging>=20.0 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (23.0)
Requirement already satisfied: pillow>=6.2.0 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from matplotlib==3.8.0) (2.8.2)
Requirement already satisfied: six>=1.5 in /Applications/Anaconda/anaconda3/lib/python3.11/site-packages (from python-dateutil>=2.7->matplotlib==3.8.0) (1.16.0)

Using cached matplotlib-3.8.0-cp311-cp311-macosx_11_0_arm64.whl (7.5 MB)

Installing collected packages: matplotlib

Attempting uninstall: matplotlib

Found existing installation: matplotlib 3.7.1

Uninstalling matplotlib-3.7.1:

Successfully uninstalled matplotlib-3.7.1

Successfully installed matplotlib-3.8.0

Note: you may need to restart the kernel to use updated packages.

Collecting numpy==1.26.1

Obtaining dependency information for numpy==1.26.1 from https://files.pythonhosted.org/packages/e8/06/0512e2582fd27bb7b358fa1e4ffc0f6c89c89f5ada31df58c5fa93171098/numpy-1.26.1-cp311-cp311-macosx_11_0_arm64.whl.metadata

Downloading numpy-1.26.1-cp311-cp311-macosx_11_0_arm64.whl.metadata (115 kB)
115.1/115.1 kB 3.6 MB/s eta 0:00:

00

Using cached numpy-1.26.1-cp311-cp311-macosx_11_0_arm64.whl (14.0 MB)

Installing collected packages: numpy

Attempting uninstall: numpy

Found existing installation: numpy 1.24.3

Uninstalling numpy-1.24.3:

Successfully uninstalled numpy-1.24.3

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

tables 3.8.0 requires blosc2~=2.0.0, which is not installed.

tables 3.8.0 requires cython>=0.29.21, which is not installed.

gensim 4.3.0 requires FuzzyTM>=0.4.0, which is not installed.

numba 0.57.0 requires numpy<1.25,>=1.21, but you have numpy 1.26.1 which is incompatible.

tensorflow-macos 2.13.0 requires numpy<=1.24.3,>=1.22, but you have numpy 1.26.1 which is incompatible.

Successfully installed numpy-1.26.1

Note: you may need to restart the kernel to use updated packages.

Packages

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

pd.options.display.max_columns = None
pd.options.display.max_rows = None
%matplotlib inline
```

Read the file

```
In [4]: df = pd.read_csv('./cereal.csv')
df.head()
```

```
Out[4]:
```

| | name | mfr | type | calories | protein | fat | sodium | fiber | carbo | sugars | potass | vitamins | sl |
|---|---------------------------|-----|------|----------|---------|-----|--------|-------|-------|--------|--------|----------|----|
| 0 | 100% Bran | N | C | 70 | 4 | 1 | 130 | 10.0 | 5.0 | 6 | 280 | 25 | |
| 1 | 100% Natural Bran | Q | C | 120 | 3 | 5 | 15 | 2.0 | 8.0 | 8 | 135 | 0 | |
| 2 | All-Bran | K | C | 70 | 4 | 1 | 260 | 9.0 | 7.0 | 5 | 320 | 25 | |
| 3 | All-Bran with Extra Fiber | K | C | 50 | 4 | 0 | 140 | 14.0 | 8.0 | 0 | 330 | 25 | |
| 4 | Almond Delight | R | C | 110 | 2 | 2 | 200 | 1.0 | 14.0 | 8 | -1 | 25 | |

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   name        77 non-null    object
1   mfr         77 non-null    object
2   type        77 non-null    object
3   calories    77 non-null    int64
4   protein     77 non-null    int64
5   fat         77 non-null    int64
6   sodium      77 non-null    int64
7   fiber       77 non-null    float64
8   carbo       77 non-null    float64
9   sugars      77 non-null    int64
10  potass      77 non-null    int64
11  vitamins    77 non-null    int64
12  shelf       77 non-null    int64
13  weight      77 non-null    float64
14  cups        77 non-null    float64
15  rating      77 non-null    float64
dtypes: float64(5), int64(8), object(3)
memory usage: 9.8+ KB
```

In [6]: `# Chek for null values`
`df.isna().sum().sum()`

Out[6]: 0

Questions

Question 1. (4 points)

- Create a new 'Type of Cereal' column in your dataframe (1 point) by copying the 'name' column. Write a function to replace the names of the cereal in your new column with one of these categories Bran, Wheat, Fiber, Protein, Crunch, Corn, Nut, Rice and Other (3 points). Hint: the function should look through the text in the cereal name and determine, based on its contents, how to categorize the cereal type.

In [7]: `# Copy the column 'name' to create the new column 'Type of Cereal'`
`df['Type of Cereal'] = df['name'].copy(deep=True)`
`df.head()`

Out[7]:

| | name | mfr | type | calories | protein | fat | sodium | fiber | carbo | sugars | potass | vitamins | sl |
|---|---------------------------|-----|------|----------|---------|-----|--------|-------|-------|--------|--------|----------|----|
| 0 | 100% Bran | N | C | 70 | 4 | 1 | 130 | 10.0 | 5.0 | 6 | 280 | 25 | |
| 1 | 100% Natural Bran | Q | C | 120 | 3 | 5 | 15 | 2.0 | 8.0 | 8 | 135 | 0 | |
| 2 | All-Bran | K | C | 70 | 4 | 1 | 260 | 9.0 | 7.0 | 5 | 320 | 25 | |
| 3 | All-Bran with Extra Fiber | K | C | 50 | 4 | 0 | 140 | 14.0 | 8.0 | 0 | 330 | 25 | |
| 4 | Almond Delight | R | C | 110 | 2 | 2 | 200 | 1.0 | 14.0 | 8 | -1 | 25 | |

```
In [8]: def replace_cereal_type():
        """
        Replace value of 'Type of Cereal' column base in its name
        """

        # Create a list of categories and replace the values of the cell using regular
        catetories = ['Bran', 'Wheat', 'Fiber', 'Protein', 'Crunch', 'Corn', 'Nut', 'Rice']
        for category in catetories:
            df['Type of Cereal'].replace(regex=fr'.*{category}.*', value=category,
                                         inplace=True)

        # Replace the value for 'Other' if the value is not in a category.
        df['Type of Cereal'].replace(regex=r'^(?!Bran|Wheat|Fiber|Protein|Crunch|Corn|Nut|Rice)',
                                     value='Other', inplace=True)

        replace_cereal_type()
```

```
In [9]: # Check the unique values of the column Type of Cereal
        df['Type of Cereal'].unique()
```

```
Out[9]: array(['Bran', 'Other', 'Crunch', 'Corn', 'Wheat', 'Nut', 'Rice'],
              dtype=object)
```

Question 2. (2 points)

- Identify the negative values in the data set and replace them with the median value for that column.

```
In [10]: # Get the min value of the columns to identify if they have negative values
        df.min()
```

```
Out[10]: name          100% Bran
mfr          A
type         C
calories     50
protein      1
fat          0
sodium       0
fiber        0.0
carbo        -1.0
sugars       -1
potass       -1
vitamins     0
shelf        1
weight       0.5
cups         0.25
rating       18.042851
Type of Cereal Bran
dtype: object
```

The dataset shows negative values in the columns: carbo, sugars, potass

```
In [11]: def replace_negative_values(df, column_name):
        """
        Replace nevasive values for the median
        ...
        Arguments
            df: DataFramw to process
            column_name: Column name
        """

        # Filter items with negative values
        filter = df[column_name] < 0
        count = len(df[filter])

        # Get the median for the column and replace values
        median = df[column_name].median()
        df.loc[filter, column_name] = median

        print(f"– Replacing {count} values on '{column_name}' column. Median = {median}")

        # Get numeric columns
        negative_columns = ['carbo', 'sugars', 'potass']
        for column_name in negative_columns:
            replace_negative_values(df, column_name)

        – Replacing 1 values on 'carbo' column. Median = 14.0
        – Replacing 1 values on 'sugars' column. Median = 7.0
        – Replacing 2 values on 'potass' column. Median = 90.0
```

```
In [12]: # Get the mean value of the columns to validate nevasive values again
df[negative_columns].min()
```

```
Out[12]: carbo      5.0
sugars    0.0
potass    15.0
dtype: float64
```

Now, there are no negative values in the columns.

Question 3. (5 points)

- Standardize the 'weight' column to 1. For this question, you will need to write a function to divide the remaining columns which contain nutritional information by the corresponding value in the weight column, and you will need to divide the value in the weight column by itself in order to get 1. For example, if an observation has a weight value of 1.33 and a calories value of 250, if you divide 250/1.33 you should get a calories value of 188 and a weight value of 1.

```
In [13]: def standarize_values():
        """
        Standarize the values of the columns base on the 'weight' column
        """
        df.loc[:, 'calories': 'weight'] = df.loc[:, 'calories': 'weight'].div(df['weight'])

    standarize_values()
    df.head(10)
```

```
Out[13]:
```

| | name | mfr | type | calories | protein | fat | sodium | fiber | carbo |
|---|---------------------------|-----|------|------------|----------|----------|------------|-----------|-----------|
| 0 | 100% Bran | N | C | 70.000000 | 4.000000 | 1.000000 | 130.000000 | 10.000000 | 5.000000 |
| 1 | 100% Natural Bran | Q | C | 120.000000 | 3.000000 | 5.000000 | 15.000000 | 2.000000 | 8.000000 |
| 2 | All-Bran | K | C | 70.000000 | 4.000000 | 1.000000 | 260.000000 | 9.000000 | 7.000000 |
| 3 | All-Bran with Extra Fiber | K | C | 50.000000 | 4.000000 | 0.000000 | 140.000000 | 14.000000 | 8.000000 |
| 4 | Almond Delight | R | C | 110.000000 | 2.000000 | 2.000000 | 200.000000 | 1.000000 | 14.000000 |
| 5 | Apple Cinnamon Cheerios | G | C | 110.000000 | 2.000000 | 2.000000 | 180.000000 | 1.500000 | 10.500000 |
| 6 | Apple Jacks | K | C | 110.000000 | 2.000000 | 0.000000 | 125.000000 | 1.000000 | 11.000000 |
| 7 | Basic 4 | G | C | 97.744361 | 2.255639 | 1.503759 | 157.894737 | 1.503759 | 13.533835 |
| 8 | Bran Chex | R | C | 90.000000 | 2.000000 | 1.000000 | 200.000000 | 4.000000 | 15.000000 |
| 9 | Bran Flakes | P | C | 90.000000 | 3.000000 | 0.000000 | 210.000000 | 5.000000 | 13.000000 |

Question 4. (5 points)

- Create a new column to categorize cereals as 'healthy' vs. 'unhealthy'. You can define your own version of healthy vs. unhealthy, or you can use the following: healthy cereals can be defined as those which have low calories (<100), low sodium (<150), low sugar (<9) high fiber (>3), and high protein (>2). All other cereals are unhealthy.


```
In [14]: def categorize_cereals():
        """
        This function categorizes dataset rows into 'healthy' vs. 'unhealthy' based on
        """
        max_calories = 100 # 100
        max_sodium = 150 # 150
        max_sugar = 9 # 9
        min_fiber = 3 # 3
        min_protein = 2 # 2

        filter = (df['calories'] < max_calories) & (df['sodium'] < max_sodium) & (
        df['is_healthy'] = np.where(filter, 'healthy', 'unhealthy')

        # Get 3 healthy and 3 unhealthy items (Only for visualization)
        df.groupby(['is_healthy'], group_keys=False).apply(lambda x: x.sample(3))

        categorize_cereals()
```

```
In [15]: df.head()
```

```
Out[15]:
```

| | name | mfr | type | calories | protein | fat | sodium | fiber | carbo | sugars | potass | vitamins | sl |
|---|---------------------------|-----|------|----------|---------|-----|--------|-------|-------|--------|--------|----------|----|
| 0 | 100% Bran | N | C | 70.0 | 4.0 | 1.0 | 130.0 | 10.0 | 5.0 | 6.0 | 280.0 | 25.0 | |
| 1 | 100% Natural Bran | Q | C | 120.0 | 3.0 | 5.0 | 15.0 | 2.0 | 8.0 | 8.0 | 135.0 | 0.0 | |
| 2 | All-Bran | K | C | 70.0 | 4.0 | 1.0 | 260.0 | 9.0 | 7.0 | 5.0 | 320.0 | 25.0 | |
| 3 | All-Bran with Extra Fiber | K | C | 50.0 | 4.0 | 0.0 | 140.0 | 14.0 | 8.0 | 0.0 | 330.0 | 25.0 | |
| 4 | Almond Delight | R | C | 110.0 | 2.0 | 2.0 | 200.0 | 1.0 | 14.0 | 8.0 | 90.0 | 25.0 | |

Question 5. (2 points)

- Based on your newly prepared data set, identify what % of cereals that each manufacturer produces are healthy.

```
In [16]: # List the manufacturer
df['mfr'].unique()
```

```
Out[16]: array(['N', 'Q', 'K', 'R', 'G', 'P', 'A'], dtype=object)
```

```
In [17]: def calculate_healthy_percentage_q5() -> pd.DataFrame:
        """
        Calculates the percentage of healthy cereals produced by the manufacturer and
        new DataFrame with the name of the manufacturer and the percentage sorted by
        """
        # Group rows by on Manufacturer and is_healthy
```

```

group = df[['mfr', 'is_healthy']].groupby(['mfr'])
result_df = group['is_healthy'].apply(lambda x: (x == 'healthy').sum()) / (
return result_df.rename('%').reset_index()

question5_df = calculate_healthy_percentage_q5()
question5_df.sort_values(by='%', ascending=False)

```

Out[17]:

| | mfr | % |
|---|-----|----------|
| 3 | N | 0.500000 |
| 4 | P | 0.111111 |
| 2 | K | 0.043478 |
| 0 | A | 0.000000 |
| 1 | G | 0.000000 |
| 5 | Q | 0.000000 |
| 6 | R | 0.000000 |

Question 6. (2 points)

- Calculate the average, minimum and maximum ratings for healthy vs. unhealthy cereals.

In [18]:

```

def calculate_average_q6(input_df) -> pd.DataFrame:
    """
    Calculate the average, minimum and maximum ratings for healthy vs. unhealthy cereals.
    """
    result = input_df.groupby('is_healthy')['rating'].agg(['mean', 'max', 'min'])
    return result

question6_df = calculate_average_q6(df)
question6_df

```

Out[18]:

| | is_healthy | mean | max | min |
|---|------------|-----------|-----------|-----------|
| 0 | healthy | 69.146753 | 93.704912 | 40.917047 |
| 1 | unhealthy | 40.826743 | 72.801787 | 18.042851 |

Question 7. (2 points)

- Calculate the average, minimum and maximum ratings for each type of cereal: Bran, Wheat, Fiber, Protein, Crunch, Corn, Nut, Rice and Other.

In [19]:

```

def calculate_ratings_q7(input_df):
    """
    Calculate the average, minimum and maximum ratings for each type of cereal.
    """
    result = input_df.groupby('Type of Cereal')['rating'].agg(['mean', 'max', 'min'])
    return result

question7_df = calculate_ratings_q7(df)
question7_df

```

Out [19]:

| | Type of Cereal | mean | max | min |
|---|----------------|-----------|-----------|-----------|
| 0 | Bran | 50.714179 | 93.704912 | 28.592785 |
| 1 | Corn | 40.482720 | 45.863324 | 35.782791 |
| 2 | Crunch | 26.078598 | 36.523683 | 18.042851 |
| 3 | Nut | 42.736791 | 53.371007 | 31.072217 |
| 4 | Other | 37.379947 | 55.333142 | 21.871292 |
| 5 | Rice | 47.771735 | 60.756112 | 40.560159 |
| 6 | Wheat | 56.333863 | 72.801787 | 36.176196 |

Question 8. (3 points)

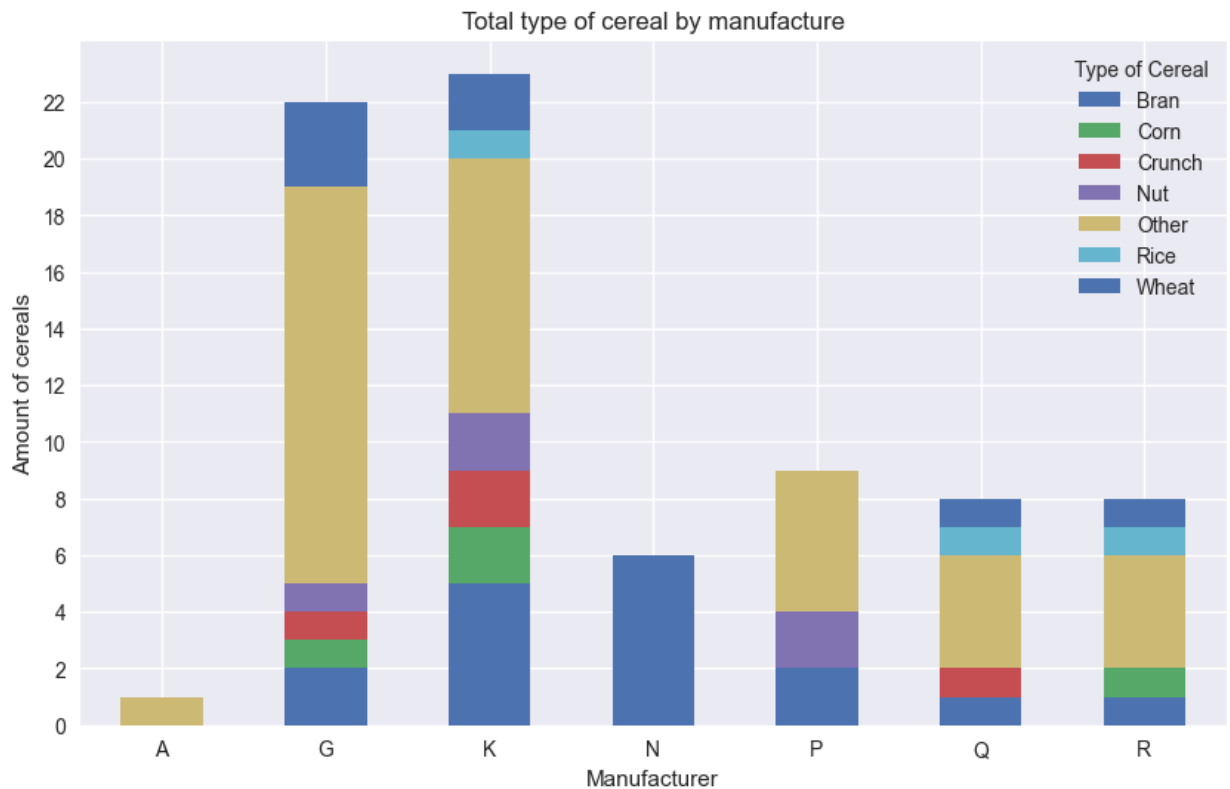
- Create a stacked bar chart which shows how many of each type of cereal each manufacturer produces.

```
In [20]: # Group rows by manufacturer
mfrgroup_df = df.groupby('mfr')['Type of Cereal'].value_counts().reset_index()
mfrgroup_df

pivot_df = mfrgroup_df.pivot(index='mfr', columns='Type of Cereal', values='count')

# Set the theme
plt.style.use('seaborn-v0_8')

# Create graph
pivot_df.plot(kind='bar', stacked=True, figsize=(10, 6))
plt.xlabel('Manufacturer')
plt.ylabel('Amount of cereals')
plt.title('Total type of cereal by manufacture')
plt.xticks(rotation=0)
plt.yticks(np.arange(0, 24, 2))
plt.show()
```



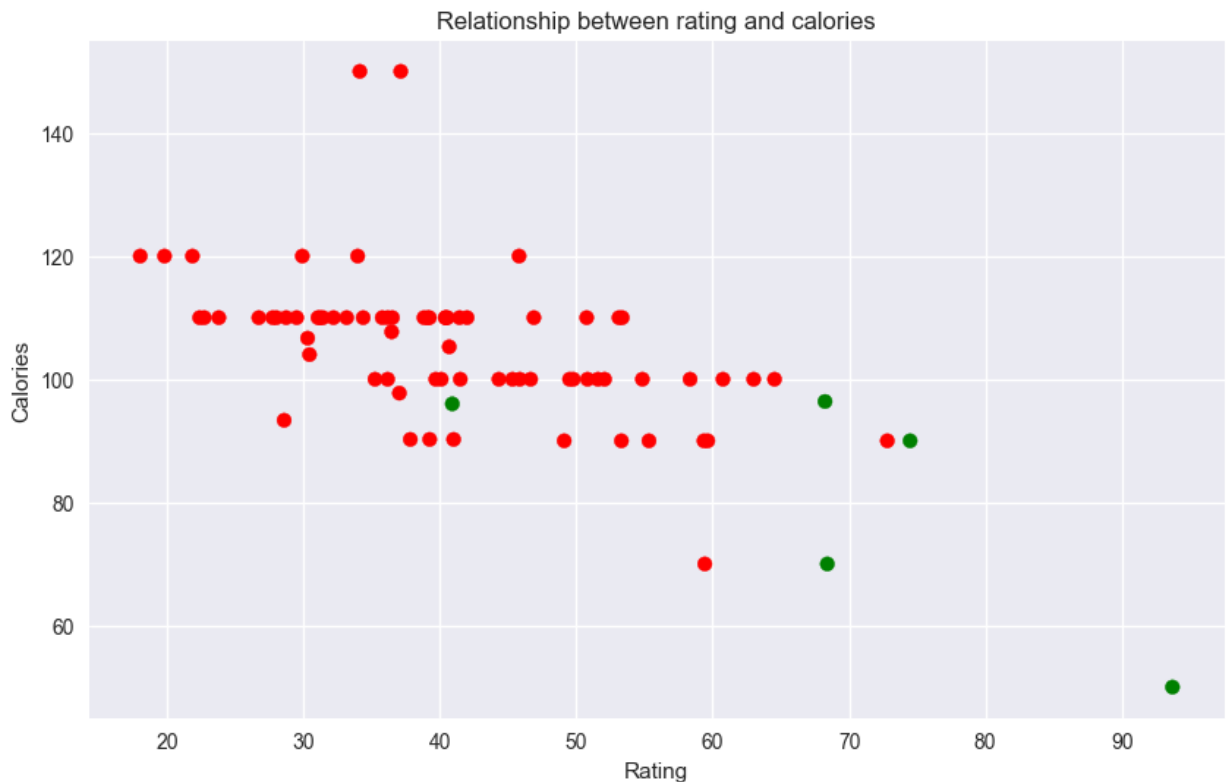
The chart shows the distribution of cereals by manufacturer. K has the highest productions and then G. A is the lower producer.

Question 9. (3 points)

- Create a 3-dimensional scatterplot which shows the relationship between rating and calories; the 3-rd dimension should be reflected in the color of the dots and should highlight whether the cereal is categorized as healthy or unhealthy.

```
In [21]: corr_df = df['is_healthy'].apply(lambda x: 'g' if x == 'healthy' else 'r')

plt.figure(figsize=(10, 6))
plt.scatter(x=df['rating'], y=df['calories'], c=corr_df)
plt.xlabel('Rating')
plt.ylabel('Calories')
plt.title("Relationship between rating and calories")
plt.show();
```



Question 10. (1 point)

- Which shelf has the most healthy cereals?

In [22]: `df[df['is_healthy'] == 'healthy'].groupby('shelf')['is_healthy'].value_counts()`

Out[22]:

| | shelf | is_healthy | count |
|---|----------|------------|-------|
| 3 | 3.000000 | healthy | 2 |
| 0 | 1.000000 | healthy | 1 |
| 1 | 1.204819 | healthy | 1 |
| 2 | 2.400000 | healthy | 1 |

The shelf with the most healthy products is 3.000000

Question 11. (1 point)

- Based on the analysis conducted, what can you conclude about the cereal data set?

Conclusion

- The dataset contains 77 rows and 16 columns, 8 int64 columns, 5 float columns and 3 string columns; it does not have any missing values, and the products were categorized by type of cereal, and it contains seven categories: Bran, Other, Crunch, Corn, Wheat, Nut and Rice.

2. The columns carbo, sugars and potass have negative values, and their values were replaced with the mean of each column because they are incorrect values.
3. Manufacturer N has the healthiest products among the companies, with 50% of healthy products, followed by P with 11% and K with 4%.
4. Healthy products have a higher rating than unhealthy products, with a maximum rating of 69.146753 against 40.826743.
5. The stacked bar chart shows the distribution of cereals by manufacturer. K has the highest production, and then G. A is the lower producer.
6. Finally, the scatter plot shows that only five products are healthy based on the thresholds.